



Part 2: Architectural Explanation

The CAF system is designed with a three-layered architecture that separates concerns effectively.

Command Processing flows top-down: The user interacts with the top-level CLI (Python) layer, where **main** serves as the entry point. `cli` handles argument parsing, and `cli_commands` orchestrates the high-level logic for each user command. This layer then delegates all control to the middle `libcaf` (Python) layer. Here, the repository module acts as the central brain, managing the state of the repository by using `ref` for reference handling and plumbing to access low-level functionality. The final layer, `src` (C++), contains the high-performance logic, and is responsible for data persistence – this is the layer where serialization of commit and tree objects is taking place, as well as file I/O operations.

The `libcaf` layer communicates with the `src` layer through the `bind` module, which exposes C++ functions and data structures to Python – making language integration possible. Lastly, error handling propagates upwards: a low-level C++ error is converted by `pybind11` into a Python exception, which is caught by the `cli_commands` layer to provide a user-friendly error message (for instance, a runtime failure to open a file can propagate upward and appear to the user as a repository error). On top of that, errors associated with user's input are being caught instantly – without further assessment at deeper layers.